

XTS mode revisited: high hopes for key scopes?

Milan Brož^{1,2}
gmazyland@gmail.com

Vladimír Sedláček²
vlada.sedlacek@mail.muni.cz

¹Cryptsetup project

²Faculty of Informatics, Masaryk University, Czechia

February 25, 2025

Abstract

This paper concisely summarizes the XTS block encryption mode for storage sector-based encryption applications and clarifies its limitations. In particular, we aim to provide a unified basis for much needed discussions about the newly proposed key scope change to the IEEE 1619 standard.

This work is licensed under a “CC BY 4.0” license.

1 Introduction

Since the introduction of XTS [3] in 2007, many software and hardware sector-based encryption systems have widely adopted it, e.g., BitLocker, VeraCrypt, Cryptsetup, and TCG Opal self-encrypting drives. Still, it seems to be surrounded by a shroud of misunderstandings, and some of its subtler points are still sometimes contested in the developer communities. It does not help that the existing NIST XTS-AES recommendation [1] only references the paid version of the IEEE standard [3], making it probably the only NIST document that does not include a publicly available definition of the described cryptographic primitive (as already mentioned in Rogaway’s analysis [11] in 2011).

The situation is especially problematic in light of the newly proposed key scope change to the standard [4] (again behind a paywall), which would render the vast majority of implementations non-compliant. This would force vendors to request (not only) open-source projects to make significant changes to adapt to the new version. In contrast, the reasons and security implications of the proposed change are not communicated at all.

The goals of this paper are threefold:

- Unify the terminology in the XTS context in a way easily understandable to different communities.
- Describe XTS security limitations and clarify some contested points.
- Encourage a constructive public discussion about future requirements and recommendations.

2 Terminology and definitions

2.1 Units and sectors

The storage industry is known to mix the power-of-10 units (SI) and the power-of-2 units for a storage size. To avoid any confusion, we strictly use powers of two with IEC prefixes: a byte is 8 bits, a kibibyte (KiB) is 2^{10} bytes, a mebibyte (MiB) is 2^{20} bytes (1024 KiB), and a tebibyte (TiB) is 2^{40} bytes (1024 MiB).

A block device is an abstraction of either a physical device (like a hard disk or a flash-based device) or a virtualized device that can be accessed (read, written to, or erased) only in units called sectors. A typical sector size is 4096 bytes (4 KiB) or 512 bytes (for older devices). In this text, we will only consider encryption of block devices (even though XTS can be used in other scenarios). Furthermore, we assume that the sector is always read or written to atomically, and individual sectors are accessed independently.¹

The typical blocksize for modern ciphers is much smaller than the sector size (16 bytes for AES). Thus, to encrypt the whole sector, we need to use an encryption mode (such as XTS) that builds upon the original cipher. All modern real-world storage devices always use sector sizes that are multiples of 16 bytes, eliminating the need for padding during both encryption and decryption.²

2.2 Notation

In this text, E will denote a blockcipher (e.g., AES-128) of blocksize n bits, so that for each key K , E_K is a permutation on $\{0, 1\}^n$ – the set of binary n -bit strings. For $s, s' \in \{0, 1\}^n$, we will write $s \| s'$ and $s \oplus s'$, for the concatenation and XOR respectively. The $s = a_{n-1} \dots a_1 a_0$ can be seen as the class of the polynomial $a_{n-1}x^{n-1} + \dots + a_1x + a_0$ in the finite field \mathbb{F}_{2^n} . With this interpretation, we can multiply binary strings by elements of \mathbb{F}_{2^n} using shift and XOR operations (though it depends on the polynomial chosen to represent \mathbb{F}_{2^n}).

S will denote the total number of sectors on the device and J the number of n -bit blocks within a sector. Correspondingly, N will always denote a sector number and j the number of the n -bit block within the sector.

We will denote the unencrypted data (plaintext) at block j within sector N by $P_{N,j}$ and naturally extend this by concatenation:

$$P_N := P_{N,0} \| P_{N,1} \| \dots \| P_{N,J-1}$$

will denote the unencrypted data at sector number N , while

$$P := P_0 \| P_1 \| \dots \| P_{S-1}$$

will denote the unencrypted data of the whole device. Analogically, we denote the encrypted data by $C_{N,j} := E_K(P_{N,j})$, C_n, C , respectively.

¹Modern storage devices are complex systems that can internally use different storage blocks for optimal performance, and can have specific requirements for optimal operations.

²The only legacy exception is the 520-byte sector (a 512-byte data sector with an additional 8-byte integrity field) on enterprise devices.

3 XTS encryption mode

The XTS mode – or rather XTS-AES – was first defined in Annex D.4.3 in [3] as an “instantiation” of Rogaway’s XEX [10] with the underlying blockcipher being AES-128 or AES-256. Though not quite, as there are several differences:

- While XEX uses a single symmetric key to encrypt both the data and the sector number, XTS uses two different symmetric keys: K to encrypt the data, and K_T to encrypt the sector number.
- While XEX starts at $j = 1$, XTS starts at $j = 0$.
- XTS allows for ciphertext stealing to deal with encrypting data whose size is not a multiple of n . As explained in Section 2.1, we focus only on full sector encryption, so we can ignore this.³

3.1 XTS definition

We propose a definition that seeks a compromise between Rogaway’s original XEX formality and the practical use case at hand.

Definition 1 (XTS encryption mode). *Let K , K_T be symmetric keys for a blockcipher E of blocksize⁴ 128 bits. Then the XTS mode of E encrypts the storage data in the following way:*

$$\begin{aligned} C_{N,j} &:= E_K(P_{N,j} \oplus T_{N,j}) \oplus T_{N,j}, \\ C_N &:= C_{N,0} \| C_{N,1} \| \dots \| C_{N,J-1}, \\ C &:= C_0 \| C_1 \| \dots \| C_{S-1}, \end{aligned}$$

where

$$T_{N,j} := E_{K_T}(N) \cdot \alpha^j$$

and $\alpha := 0^{126}10$ is the class of x in $\mathbb{F}_{2^{128}} := \mathbb{F}_2[x]/(x^{128} + x^7 + x^2 + x + 1)$.

Note that for a binary string $a_{127}a_{126} \dots a_1a_0$, we have

$$s \cdot \alpha = \begin{cases} a_{126}a_{125} \dots a_0 0 & \text{if } a_{127} = 0, \\ a_{126}a_{125} \dots a_0 0 \oplus 0^{120}10^4 1^3 & \text{if } a_{127} = 1, \end{cases}$$

as x^{128} and $x^7 + x^2 + x + 1$ represent the same element in $\mathbb{F}_{2^{128}}$.

³XTS has been used in the Linux kernel since 2007. In contrast, ciphertext stealing was only added in 2019, and probably there is still no active in-kernel user.

⁴This is not essential, but multiplication by α differs for different block sizes.

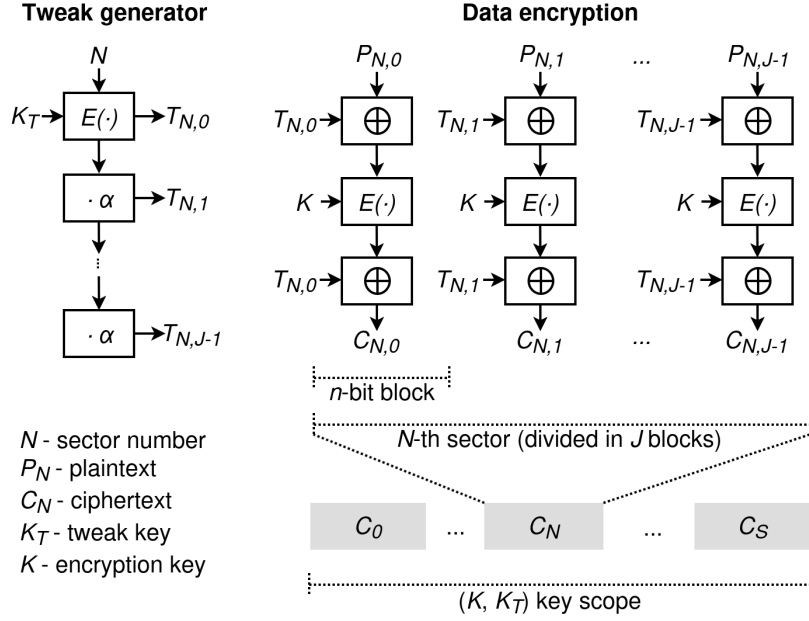


Figure 1: The principle of XTS sector encryption (without ciphertext stealing).

Example 1 (AES-XTS-128 encryption of one sector with two blocks).

For

$$\begin{aligned} N &= 0x00000000000000000000000000000001, \\ K &= 0x11111111111111111111111111111111, \\ K_T &= 0x22222222222222222222222222222222, \\ P_{N,0} &= 0x44444444444444444444444444444444, \\ P_{N,1} &= 0x88888888888888888888888888888888, \end{aligned}$$

the tweak values are generated in the following way:

$$\begin{aligned} T_{N,0} &= E_{K_T}(N) = 0x6752ca5febca0f3fc8dc9dfc2a916295, \\ T_{N,1} &= T_{N,0} \cdot \alpha^1 = 0x49a494bfd6951f7e90b93bf95522c52a. \end{aligned}$$

The sector blocks are then encrypted as follows:

$$\begin{aligned} C_{N,0} &= E_K(P_{N,0} \oplus T_{N,0}) \oplus T_{N,0} \\ &= E_K(0x6752ca5febca0f3fc8dc9dfc2a916295) \oplus T_{N,0} \\ &= 0x13f084e65a7ca361f74957c9b11c7710 \oplus T_{N,0} \\ &= 0x74a24eb9b1b6ac5e3f95ca359b8d1585 \end{aligned}$$

$$\begin{aligned} C_{N,1} &= E_K(P_{N,1} \oplus T_{N,1}) \oplus T_{N,1} \\ &= E_K(0xc12c1c375e1d97f61831b371ddaa4da2) \oplus T_{N,1} \\ &= 0x2cada9d22ad34bf19a226c2c824f0364 \oplus T_{N,1} \\ &= 0x65093d6dfc46548f0a9b57d5d76dc64e \end{aligned}$$

3.2 Threat model

The simplest threat model for length-preserving sector encryption usually considers a “stolen device”, allowing the attacker to access only one ciphertext snapshot at time. A typical example of such a threat model is the TCG Opal definition [2]: “*protect the confidentiality of stored user data against unauthorized access once it leaves the owner’s control (following a power cycle and subsequent deauthentication).*” This model expects the user to recognize when the protected device “left the owner’s control”, which is not always feasible.

The real-world threat model also considers situations where the attacker can manipulate the ciphertext after repeatable access to the encrypted device, even performing some form of “traffic analysis”. This is important, as encrypted images are often stored in cloud environments. Section D.4.3 in [4] suggests (in addition to introducing a key scope): “*The decision on the maximum amount of data to be encrypted with a single key should take into account the above calculations together with the practical implication of the described attack (e.g., ability of the adversary to modify plaintext of a specific block, where the position of this block may not be under adversary’s control).*”.

Relatedly, we should consider adversaries capable of the chosen-ciphertext attack (CCA), i.e., able to query both the encryption and decryption oracle (with K, K_T still being fixed and secret) on inputs of their choice.

Despite having received a lot of criticism [11, 9] after introduction, the XTS mode still has not been replaced by any other mode more than a decade later. Unfortunately, the security goals and threat models are not well-defined either.

4 Security limits of XTS

4.1 Key scopes

The most drastic proposed change in [4] is the following requirement: “*Maximum Number of 128-bit blocks in a key scope = 2^{36} to 2^{44}* ”. For a fixed keypair (K, K_T) , this translates to $S \cdot J \leq 2^{36}$ and $S \cdot J \leq 2^{44}$, respectively.

The reason for this comes from Annex D.4.2 in [4], which considers a CCA adversary with access to triples $(P_{N,j}, C_{N,j}, T_{N,j})$ and $(P_{N',j'}, C_{N',j'}, T_{N',j'})$ for some N, N', j, j' , such that

$$P_{N,j} \oplus T_{N,j} = P_{N',j'} \oplus T_{N',j'}. \quad (1)$$

If such an adversary can overwrite $P_{N,j}$, they can use the resulting ciphertext to change $C_{N',j'}$ to a value that decrypts to any plaintext of their choice.

There is a non-negligible probability that (1) occurs when the number of plaintexts ($S \cdot J$) encrypted under the same keypair (K, K_T) approaches the birthday bound $2^{n/2}$.

Consequently, the scope of each keypair should be limited. The probability of (1) occurring for AES- $\{128, 256\}$ can be estimated by $\frac{(S \cdot J)^2}{2^{128}}$. Thus, as argued in Annex D.4.3 in [3], encrypting 1 TiB of data (i.e., $S \cdot J = 2^{36}$) would lead to (1) occurring with probability roughly $2^{36 \cdot 2 - 128} = 2^{-53}$, whereas 256 TiB of data would correspond to roughly $2^{44 \cdot 2 - 128} = 2^{-40}$.

However, it is unclear what amount of risk is acceptable in different contexts, even though the proposed change would invalidate compliance of many current encrypted storage devices.

4.2 Maximal sector size

Section 5.1 in [3] states: “*The number of 128-bit blocks within the data unit shall not exceed 2^{20} .*” As their data unit translates to our sector, this limits its size to 16 MiB, i.e., $J \leq 2^{20}$. This condition has already been present in [3] as a change from the initial version (2007), where it was only a suggestion (“*should not exceed*”). NIST [1] also strictly mandates this limit, but the rationale has never been explained. As XEX has been proven to be secure against chosen-ciphertext attacks even without this condition [10, 6], its purpose remains unclear. In practice, this is almost never an issue, as the typical sector size is at most 4 KiB, i.e., $J \leq 2^8$ (Section 2.1).

4.3 Distinct K, K_T

Annex C.I in [7] warns about a chosen ciphertext attack against XTS-AES, referring to Section 6 in Rogaway’s original paper [10]: “*by obtaining the decryption of only one chosen ciphertext block in a given data sector, an adversary who does not know the key may be able to manipulate the ciphertext in that sector so that one or more plaintext blocks change to any desired value*”. The attack crucially relies on two simultaneous conditions: $K = K_T$, and using a tweak corresponding to a representative of 1 – i.e., in our case, $j = 0$. There are two obvious ways to prevent the attack:

- Starting from $j = 1$ (as XEX does); then the original security proof applies (Section 6 in [10], Section 4.2 in [6]).
- Requiring that $K \neq K_T$. FIPS 140-3 [7] made this option mandatory for compliance, even though as Liskov and Minematsu argue in [5], this comes from a misapplication of a security design practice (though does not lower the security in any way).

5 Discussion

The XTS encryption mode, as defined in IEEE 1619 [3], has been used for many years in many different systems. Thus, any compliance-breaking changes – such as the proposed mandatory use of key scopes in the IEEE 1619 draft [4] – could have a dramatic impact on the ecosystem. This would be amplified even more by the concept of FIPS certifications (which should impact the document [7]).

We are not necessarily arguing against the change; the issue described in Section 4.1 could definitely be a substantial one in some scenarios. However, to justify the increased complexity and potential for new implementation mistakes, it would be helpful to first have good answers to the following questions:

- What are the contexts with a weaker threat model? For example, in the “stolen” device model, where the attacker cannot actively modify the ciphertexts, adding key scopes would not improve security.
- What are the contexts with a stronger threat model? For example, if the attacker can store or extract snapshots of encrypted blocks outside the device, adding key scopes would not necessarily protect against the attack.

- Is there another way to strengthen XTS without the key scopes? One option would be using a blockcipher with larger n . As $n = 256$ would already be acceptable and NIST already plans to propose a standard for “wide AES” ([8]), this option should be included in future considerations. The problem is that the XTS standard [3] has hardcoded 128-bit blocks everywhere and would need to be updated.
- How should the keypairs used for key scopes be generated? It is unclear if all the keypairs used in key scopes must be generated using an approved random number generator, or if they can be derived from a master key with an approved key derivation algorithm. The second option would simplify key scope implementation.
- How to determine which keypair to use for which sector? And how to handle device resizing? Standardizing (several variants of) answers to these might help with complexity reduction.

We propose two partial answers to the last question:

- Implement key scopes in a linear fashion: use each keypair to encrypt the next (at most) 2^{44} plaintexts, then switch to the next keypair. A downside is that keypairs would not be utilized uniformly: if data is stored only on a small part of the device, the keys beyond the first one are never used. Also, for device resizing, keypairs would need to be added or removed dynamically.
- Implement key scopes with rotating keypairs: use $(N \bmod m)$ -th keypair to encrypt the N -th sector, where m is the total number of keypairs. This solution must define the maximal device size in advance (to not exceed the required key scope), but easily allows a dynamic resize (up to the maximal size). It also utilizes all keys more uniformly.

We believe that continuing this conversation publicly would be especially relevant not only for the hardware world, where vendors would like to reuse existing accelerated implementations. From a long-term perspective, it might be more beneficial to switch to a different (wide) encryption mode (as suggested in [11]) if length-preserving ciphertext is required. Or if the storage device provides space for authentication tags, authenticated encryption would be a strong candidate as well.

We would be grateful for any feedback or suggestions that could contribute to the improvement of this work. The source code of this document is available on GitHub⁵.

Acknowledgements

The authors were supported by the European Union under Grant Agreement No. 101087529 (CHESS – Cyber-security Excellence Hub in Estonia and South Moravia).

⁵<https://github.com/mbroz/xts-paper>

References

- [1] Morris J. Dworkin. SP 800-38E. Recommendation for Block Cipher Modes of Operation: The XTS-AES Mode for Confidentiality on Storage Devices. Technical report, National Institute of Standards & Technology, Gaithersburg, MD, United States, 2010. Available at <https://csrc.nist.gov/pubs/sp/800/38/e/final>.
- [2] Trusted Computing Group. Storage Security Subsystem Class: Opal, 2022. Available at <https://trustedcomputinggroup.org/resource/storage-work-group-storage-security-subsystem-class-opal/>.
- [3] IEEE. IEEE 1619-2018 – Standard for Cryptographic Protection of Data on Block-Oriented Storage Devices, 2018.
- [4] IEEE. IEEE 1619/D12 Draft Standard for Cryptographic Protection of Data on Block-Oriented Storage Devices, 2025.
- [5] Moses Liskov and Kazuhiko Minematsu. Comments on XTS-AES. *Comments to NIST, available from their web page*, 2008. Available at https://csrc.nist.gov/groups/ST/toolkit/BCM/documents/comments/XTS/XTS_comments-Liskov_Minematsu.pdf.
- [6] Kazuhiko Minematsu. Improved security analysis of XEX and LRW modes. In *International Workshop on Selected Areas in Cryptography*, pages 96–113. Springer, 2006.
- [7] NIST. Implementation Guidance for FIPS 140-3 and the Cryptographic Module Validation Program, 2024. Available at <https://csrc.nist.gov/csrc/media/Projects/cryptographic-module-validation-program/documents/fips%20140-3/FIPS%20140-3%20IG.pdf>.
- [8] NIST. PRE-DRAFT Call for Comments: NIST Proposes to Standardize a Wider Variant of AES, 2024. Available at <https://csrc.nist.gov/pubs/sp/800/197/iprd>.
- [9] Thomas Ptacek and Erin Ptacek. You Don’t Want XTS, 2014. Available at <https://sockpuppet.org/blog/2014/04/30/you-dont-want-xts/>.
- [10] Phillip Rogaway. Efficient instantiations of tweakable blockciphers and refinements to modes OCB and PMAC. In *International Conference on the Theory and Application of Cryptology and Information Security*, pages 16–31. Springer, 2004.
- [11] Phillip Rogaway. Evaluation of some blockcipher modes of operation, 2011. Available at <https://www.cs.ucdavis.edu/~rogaway/papers/modes-cryptrec.pdf>.